

Homework Assignment

HW2: Tiny-UNIX on the SAPC

Assigned: 2 October 2019

Due: class time 23 October 2019

1. Introduction

The objective of this assignment is to implement a tiny-UNIX with 3 services: read, write, and exit on the SAPC. The students are asked to convert the read, write functions in hw1's I/O package to system calls. In addition, they are asked to write the new service exit. The prototypes for the services are:

```
int read(int dev, char *buf, int nchar);    /* read nchar bytes into buf from dev */
int write(int dev, char *buf, int nchar);    /* write nchar bytes from buf to dev */
int exit(int exitcode);                    /* exit will return back to Tutor */
```

The testio program in hw1 can be modified to test the 3 services. Use a script file to capture the outputs.

2. Discussions

This assignment builds on your standalone I/O package from hw1. If you prefer, you can use the hw1 solution, which is included in the provided hw2 directory. In any case, fix yours, if necessary, up to spec first.

To make system calls, the user program needs to execute a trap instruction to transfer control to the kernel and the kernel's trap handler will perform the service. The x86 Linux syscall linkage is as follows:

- int 0x80 is the syscall instruction
- the syscall # is in eax
- the syscall args are in ebx (first), ecx (second), and edx (third).

Here are the files you need:

i) Shared between user and kernel:

- tty_public.h: device numbers
- syscall.h: syscall numbers (like the UNIX /usr/include/sys/syscall.h)

ii) Kernel files:

- ioconf.h, tty.h: i/o headers used only by kernel
- tsystem.h: syscall dispatch, kernel fn protos (a copy is posted)
- startup0.s: same as \$pclibsrc/startup0.s. Sets up stack, calls into startup1.c
- startup1.c: same as \$pclibsrc/startup.c, but calls your kernel initialization routine instead of main.
- tunix.c: has kernel init routine. It needs to call ioinit, call set_trap_gate(0x80,&syscall), and possibly other inits. tunix.c also has the code for syscallc, sys call exit, and set_trap_gate. The code for set_trap_gate is just that of set_intr_gate with the line:

Homework Assignment

```
desc->flags = GATE_P|GATE_DPL_KERNEL|GATE_INTGATE;  
Replaced by  
desc->flags = GATE_P|GATE_DPL_KERNEL|GATE_TRAPGATE;
```

You can find the `locate_idt` function in `$pcsrc/cpureg.S`. Make sure you include the following prototype if you want to call it using C:

```
extern void locate_idt(unsigned int *limitp, char ** idtp);
```

`syscallc`: first write it with a big switch statement over the various different syscalls. If you have time, upgrade it to use a sysent dispatch table.

`sysentry.s`: Trap handler's assembler envelope routine `_syscall` for the trap resulting from a system call--needs to push `eax`, `ebx`, `ecx`, `edx` on stack, where they can be accessed from C, call `_syscallc`, then pops, `iret`. Use `$pcsrc/irq4.s` as an example and modify it to fit your needs.

`io.c`: rename "read" to "sysread", etc. to avoid linking problems, since "read" is a now user-level call.

`ioconf.c`, `ioconf.h`: from `hw1`.

`tty.c`, `tty.h`, `tty_public.h`: `tty` driver from `hw1`, unchanged

iii) User-level files:

`tunistd.h`: prototypes for user mode system calls (like UNIX `/usr/include/unistd.h`)

`uprog.c`: has `main()`. Easily extended to multiple user files, or user assembler sources, as long as they follow the syscall rules and have a `_main` entry point. First example is

```
main() { write(TTY1,"hi!\n",4); }.
```

Work back to `testio.c` from `hw1`.

`ulib.s`: library set-ups for syscalls: `_read`, `_write`, `_exit`. Provided for `write`, you add `read` and `exit`.

`crt0.s`: user-level "C startup module" sets up stack, calls `_main`, does `exit` syscall. Entry point `_ustart`.
Edit `$pcsrc/startup0.s`

iv) hw1 solution files not directly used in hw2:

`io_public.h`: like `tunistd.h` above, but also lists `ioinit()`, and not `exit()`.

`testio.c`: remove `ioinit()` call here to turn into proper user program. (and note that the `kprintf`'s are only for debugging)

v) Make file

`makefile`: The provided `makefile` can make a `hw2` system by "make `U=test1`" to use `test1.c` as a user program. The default user program is `uprog.c` and it can be build by entering "make".

In the `hw2` directory, empty files (file with length =0) are provided and you can use them to test the `makefile`. (Empty files are valid `.c` and `.s` programs and `make` treats them as regular files). If you try "make `U=test1`" with the provided files, you should see compiles followed by a load with an error as follows:

Homework Assignment

```
...
io.opc: In function `write':
/home/eoneil/444/hw2/io.c:50: multiple definition of `write'
ulib.opc(.text+0x0): first defined here
....
```

This happens because both `ulib.s` and `io.c` define global symbols named "write". You need to change `write` to `syswrite` in `io.c` to fix this. We want to use "write" for the user-level system call, so the kernel needs another name for its function. `syswrite` is the Linux kernel name for its write-implementing function, so let's adopt that name.

3) The Finished Program

The idea here is that each user program to be run on the SAPC has to be separately built with `tunix`, downloaded and run. `Startup0` executes first, and transfers control to the kernel initialization in `tunix.c`, which sets up the system and starts the user code at `ustart` (calling `ustart` will do the trick). The C user startup module reinitializes the stack and calls `main`. The syscalls in the user code (in `ulib.s`, called from `test1.c`) cause execution of the system call handler in `tunix.c` (and functions called from there), returning to the user code in `ulib.s` at the `iret`. Finally the user does a syscall `exit`. The kernel gets control, and finishes up.

4) Suggested Steps

There are lots of little pieces to this system. Here is a suggested sequence to follow:

1. Get a system built. Change `io.c`'s `write` to `syswrite`, `read` to `sysread`, fill out `startup0` and `startup1`, write a tiny `tunix.c` init function that calls `ioinit`, then calls `main` (cheating for now--later it should call `ustart`), and then returns to startup, shutting down. At this point, `test1.c` just has a `main` that `kprintf`'s a message. Now it should build and run, but does no syscalls.
2. `ulib.s` is set up for `write` already, so do the `write` syscall first. Set the trap vector up in kernel init in `tunix.c`, and write `sysentry.s`—have it push registers on the stack and then call into `tunix.c`. In `tunix.c`, access the pushed registers (themselves syscall arguments from the user) via `args` to the C function, and call `syswrite`. Make `test1.c` do a simple `write`. We'll go over the trick about the syscall `args` in class.
3. Next implement syscall `exit` and add an `exit` to `test1.c`.
4. Write the proper user startup module `crt0.s` to reinitialize the stack and call into `main`, then when that returns, it does an `exit` syscall. It has entry point `_ustart`. Change the call to `main` in kernel initialization to call `ustart` now. Try a user program without its own `exit` syscall.